ISLR Chapter 4, Ex 13
7/20/15
Tatiana Romanchishina

13. Using the Boston data set, fit classification models in order to predict whether a given suburb has a crime rate above or below the median. Explore logistic regression, LDA, and KNN models using various subsets of the predictors. Describe your findings.

First I added a new binary variable that indicates whether the crime rate is above or below the median. Then I standardized the other predictors, so that their means are 0 and standard deviation is 1. Then I assigned 100 observations to a test set and the remaining 406 to the training set. (see the R script)

Now I am ready to fit different models.

- **Logistic Regression**

First I try all the predictors:

```
================================================
                        Dependent variable:
                        ---------------------------
                                crimbin
------------------------------------------------
zn                      -1.328*
                                (0.719)

indus                   0.208
                                (0.354)

chas                    0.040
                                (0.191)

nox                     4.344***
                                (0.887)

rm                      -0.102
                                (0.585)

age                     0.425
                                (0.390)

dis                     1.374**
                                (0.542)
```

rad	6.542***
	(1.636)
tax	-0.971**
	(0.450)
ptratio	0.021
	(0.319)
black	-0.454
	(0.513)
lstat	0.517
	(0.435)
medv	1.291*
	(0.692)
Constant	3.020***
	(0.817)
---	---
Observations	406
Log Likelihood	-83.782
Akaike Inf. Crit.	195.564
==========	==========
Note:	*p<0.1; **p<0.05; ***p<0.01

It looks like nox, dis, rad, and tax have a significant relationship with the crime rate.

Now I am going to try to make predictions about my test set but I will try several thresholds:
(lr.pred are the predictions of the model with the specified threshold)

- ➔ probability of crime rate being above median = 50%:

	testCrim	
lr.pred	0	1
0	73	21
1	5	1

with the error rate = 0.15.
Even though it looks like the error rate is low, it only predicts well if a suburb has a crime rate below median, but it does not predict correctly otherwise.

➔ probability of crime rate being above median = 75%:

	testCrim	
lr.pred	0	1
0	76	22
1	2	0

with the error rate = 0.24.
The error rate is even higher than before and now it cannot predict the suburbs with crime rate above median at all.

➔ probability of crime rate being above median = 25%

	testCrim	
lr.pred	0	1
0	66	3
1	12	19

with the error rate = 0.15.
The error rate is as low as in the first try, but the predictions are looking better, because now it predicts most of both classes.

- **LDA**

1) All predictors:

Call: lda(crimbin ~ ., data = trainBoston)
Prior probabilities of groups:

0	1
0.4310345	0.5689655

Group means:
… (omitted)

Coefficients of linear discriminants:

	LD1
zn	-0.28143254
indus	0.30673534
chas	-0.03874451
nox	0.70535833
rm	0.10665269
age	0.25878572
dis	0.16752656
rad	0.94919860
tax	-0.19917231
ptratio	-0.26897926
black	-0.03284464
lstat	0.12984063
medv	0.32316451

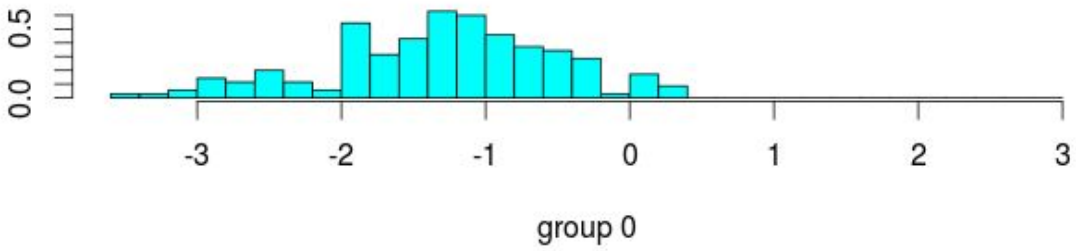

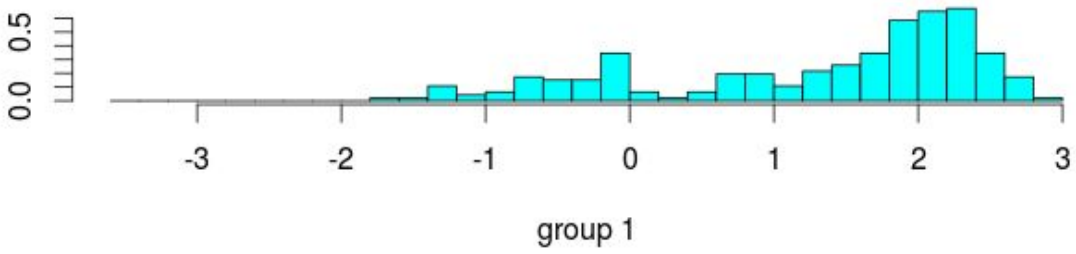

Now I can try to predict the test class using this model:

lda.class	0	1
0	76	22
1	2	0

with the error rate = 0.24. Though it is a low error rate, it does not predict the suburbs with a crime rate above the median.

2) Some of the predictors: nox, dis, rad, tax, ptratio, medv

Call: lda(crimbin ~ nox + dis + rad + tax + ptratio + medv, data = trainBoston)

Prior probabilities of groups:

0	1
0.4310345	0.5689655

Group means:

	nox	dis	rad	tax	ptratio	medv
0	-0.6515262	0.5396006	-0.6084553	-0.5520955	-0.2803591	0.3142349
1	0.8050526	-0.6929524	0.7388406	0.7232653	0.1655096	-0.2275223

Coefficients of linear discriminants:

	LD1
nox	0.89150541
dis	-0.27457496

rad 0.83731211
tax -0.06205007
ptratio -0.14532932
medv 0.22763105

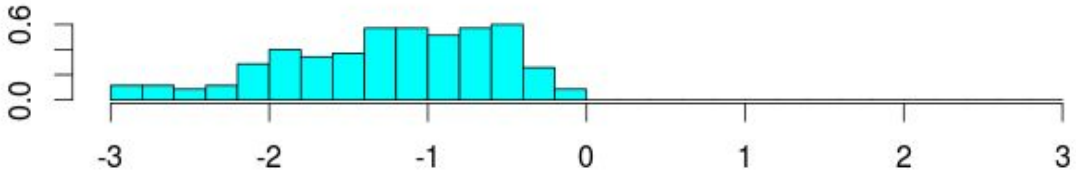

group 0

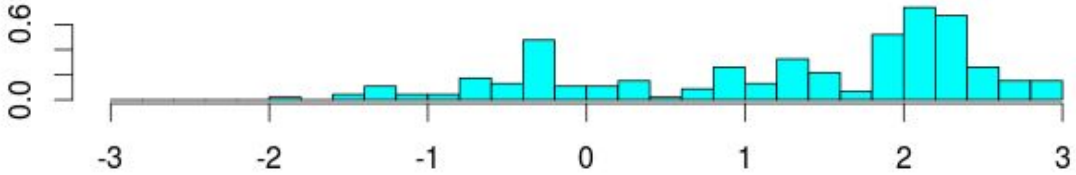

group 1

Now I can try to predict whether the crime rate is above the median or below:

	testCrim	
lda.class2	0	1
0	78	22
1	0	0

with the error rate = 0.22. Again the model does not predict correctly the suburbs that have a crime rate above the median, but it predicted correctly the suburbs that have a crime rate below the median.

- **KNN**
    - **All predictors**

Amazingly KNN model with all predictors and k=1 predicted correctly almost all instances of both classes:

	testCrim	
knn.pred	0	1
0	77	0
1	1	22

with the accuracy = 98%.

Just to confirm whether this is random or not, I am going to use another test and train sets, k=1. The results are less exciting, but still high accuracy:

	testCrim2	
knn.pred2	0	1
0	51	27
1	4	19

with the accuracy = 69%.

Now I will try using the same sets for training and testing, but try using different values of k:

- ➔ k = 3

	testCrim2	
knn.pred2	0	1
0	53	17
1	2	29

with the accuracy = 81%

- ➔ k = 10

	testCrim2	
knn.pred2	0	1
0	49	11
1	6	35

with the accuracy = 83%

- ➔ k = 20

	testCrim2	
knn.pred2	0	1
0	47	1
1	8	45

with the accuracy = 91%

It appears that increasing the value of k increases the accuracy of the model. However, the first result must have been due to “lucky” sampling. Though as one can see using a different set of training/testing data did not make the result much worse.

The accuracy of the KNN model appears to be much higher than the accuracy of the previous two models.