


5. We have seen that we can fit an SVM with a non-linear kernel in order to perform classification using a nonlinear decision boundary. We will now see that we can also obtain a nonlinear decision boundary by performing logistic regression using nonlinear transformations of the features.

(a) Generate a data set with $n = 500$ and $p = 2$, such that the observations belong to two classes with a quadratic decision boundary between them. For instance, you can do this as follows:

```
> x1 = runif(500) - 0.5  
> x2 = runif(500) - 0.5  
> y = 1*(x1^2 - x2^2 > 0)
```

(b) Plot the observations, colored according to their class labels. Your plot should display X_1 on the x-axis, and X_2 on the y-axis.



(c) Fit a logistic regression model to the data, using X_1 and X_2 as predictors.

Call:

```
glm(formula = y ~ x1 + x2, family = "binomial")
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.179	-1.139	-1.112	1.206	1.257

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.087260	0.089579	-0.974	0.330
x1	0.196199	0.316864	0.619	0.536
x2	-0.002854	0.305712	-0.009	0.993

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 692.18 on 499 degrees of freedom

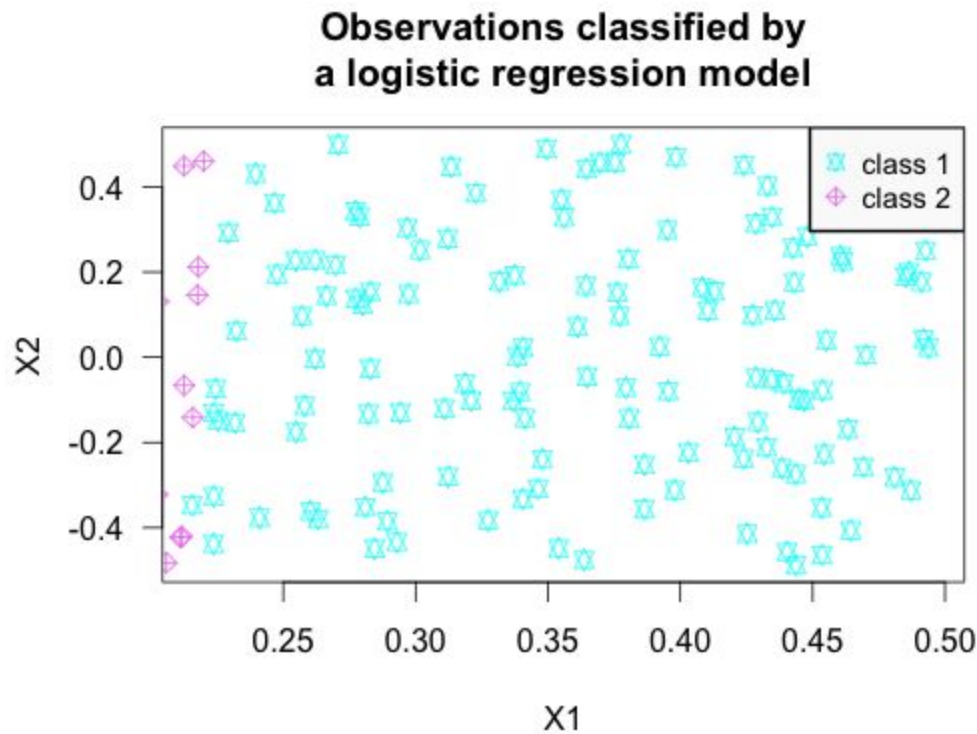
Residual deviance: 691.79 on 497 degrees of freedom

AIC: 697.79

Number of Fisher Scoring iterations: 3

The variables do not appear to have a significant relationship with the response variable.

(d) Apply this model to the training data in order to obtain a predicted class label for each training observation. Plot the observations, colored according to the predicted class labels. The decision boundary should be linear.



I used the probability threshold of .489 to produce two classes. The decision boundary seems linear and clear.

(e) Now fit a logistic regression model to the data using nonlinear functions of X1 and X2 as predictors (e.g. $X1^2$, $X1 \times X2$, $\log(X2)$, and so forth).

Call:

```
glm(formula = y ~ log2(x1) + poly(x2, 2) + I(x1 * x2), family = "binomial")
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.98542	-0.06837	0.00753	0.09873	2.37857

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	10.1540	1.9023	5.338	9.41e-08 ***
log2(x1)	4.8614	0.9048	5.373	7.75e-08 ***
poly(x2, 2)1	-41.0661	24.2062	-1.697	0.0898 .
poly(x2, 2)2	-104.6153	18.6898	-5.597	2.18e-08 ***
I(x1 * x2)	18.8225	10.5556	1.783	0.0746 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 318.830 on 229 degrees of freedom

Residual deviance: 65.559 on 225 degrees of freedom

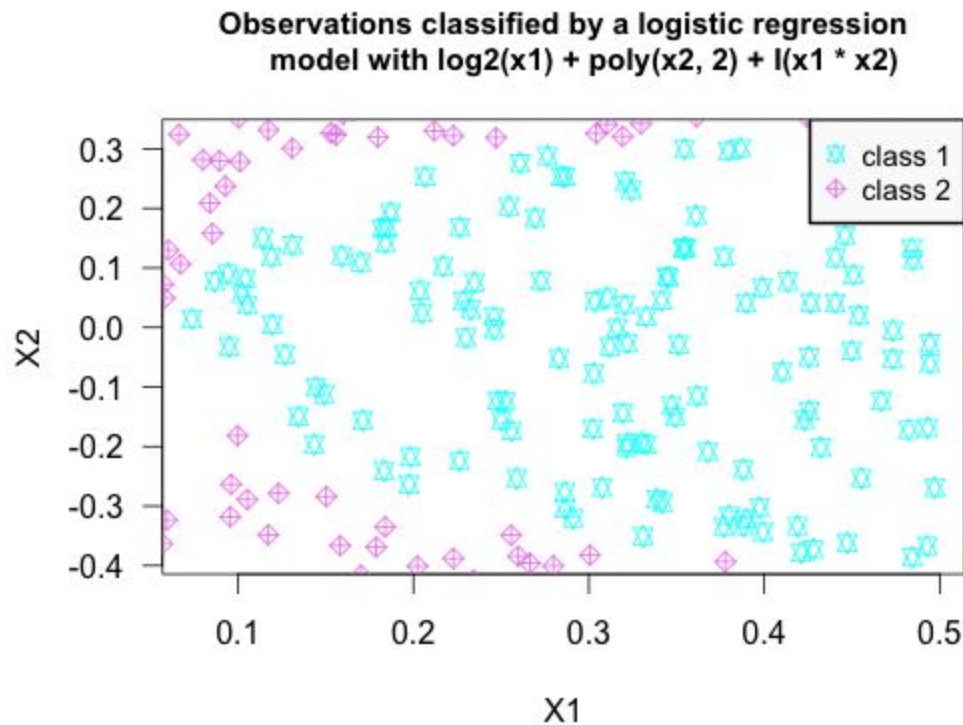
(270 observations deleted due to missingness)

AIC: 75.559

Number of Fisher Scoring iterations: 8

I used $\log2(x1) + \text{poly}(x2, 2) + I(x1 * x2)$ as independent variables, and all of the relationships appear very significant.

(f) Apply this model to the training data in order to obtain a predicted class label for each training observation. Plot the observations, colored according to the predicted class labels. The decision boundary should be obviously non-linear. If it is not, then repeat (a)-(e) until you come up with an example in which the predicted class labels are obviously non-linear.



The boundary is nonlinear, but does not resemble the original decision boundary.

(g) Fit a support vector classifier to the data with X1 and X2 as predictors. Obtain a class prediction for each training observation. Plot the observations, colored according to the predicted class labels.

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost

0.1

- best performance: 0.478

- Detailed performance results:

cost error dispersion

1 1e-05 0.484 0.04599517

2 1e-04 0.484 0.04599517

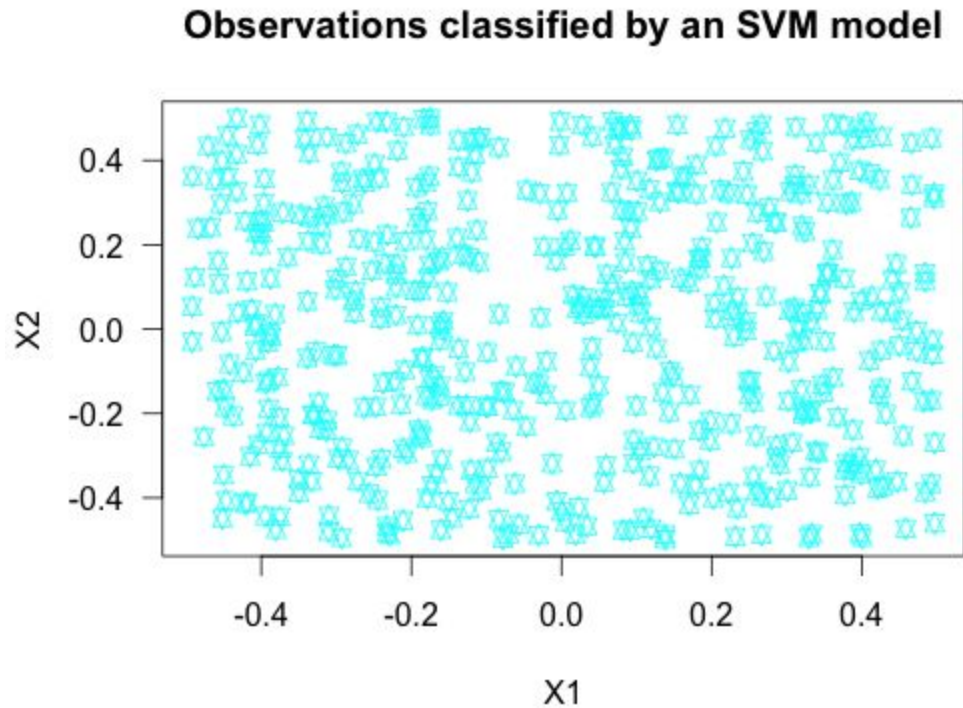
3 1e-03 0.484 0.04599517

4 1e-02 0.484 0.04599517

5 1e-01 0.478 0.05996295

6 1e+00 0.492 0.03794733
7 1e+01 0.492 0.03794733
8 5e+01 0.492 0.03794733
9 1e+02 0.492 0.03794733

Best performance is close to other results and is a pretty high error rate. All the observations are classified into one and the same class.



(h) Fit a SVM using a non-linear kernel to the data. Obtain a class prediction for each training observation. Plot the observations, colored according to the predicted class labels.

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost

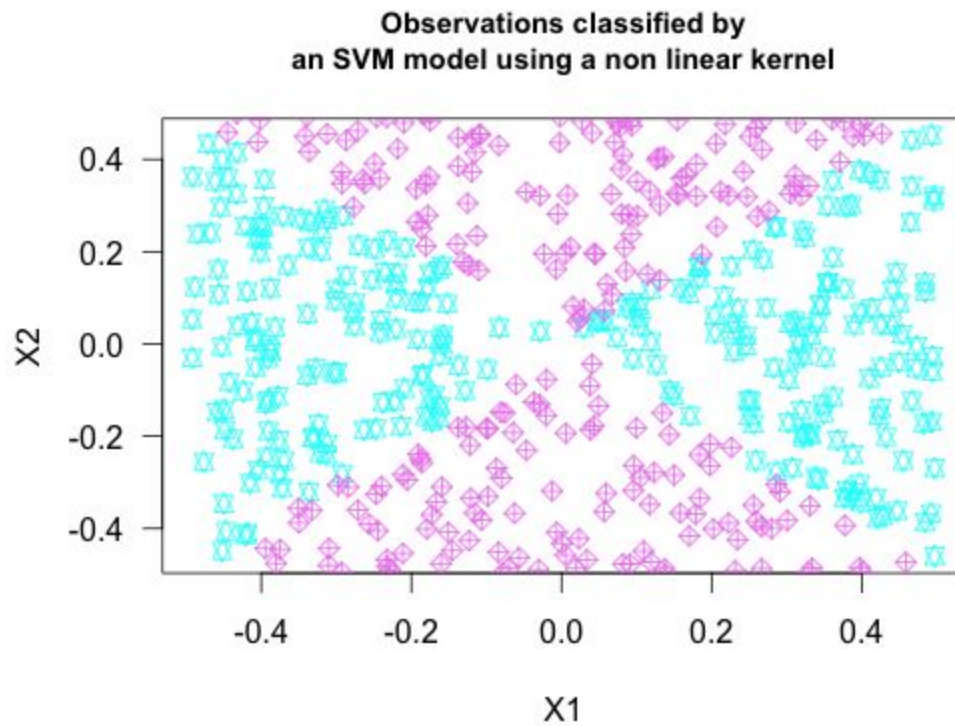
50

- best performance: 0.022

- Detailed performance results:

cost error dispersion

1	1e-05	0.456	0.15770577
2	1e-04	0.456	0.15770577
3	1e-03	0.456	0.15770577
4	1e-02	0.456	0.15770577
5	1e-01	0.084	0.03502380
6	1e+00	0.056	0.02270585
7	1e+01	0.026	0.02674987
8	5e+01	0.022	0.01988858
9	1e+02	0.024	0.02065591



(i) Comment on your results.

The SVM model with non-linear kernel produced the best result. With the help of the tune() function, I was able to quickly find out the best tuning parameter from a range of parameters.